

CS4990 – GPU Computing


Module 7: Convolution

Hao Ji


Computer Science Department

Cal Poly Pomona


- Convolution
- Parallel Convolution: a basic algorithm
- Constant Memory and Caching
- Tiled Convolution



input



filter




output

Every **output element** is a weighted sum of the neighboring **input elements**


Image blur seen before was a special case where all weights are the same

In general, weights are determined by a convolution **filter**
(commonly called convolution *kernel*, but we will use *filter* to avoid confusion with CUDA kernels)


- Commonly used in signal processing, image processing, video processing, etc.
- Usually used to transform signals or pixels to more desirable values
 - e.g., Gaussian blur, sharpen, edge detection, etc.
 - Transformation depends on the weights in the filter
- Using 2D as an example, but can also be 1D or 3D



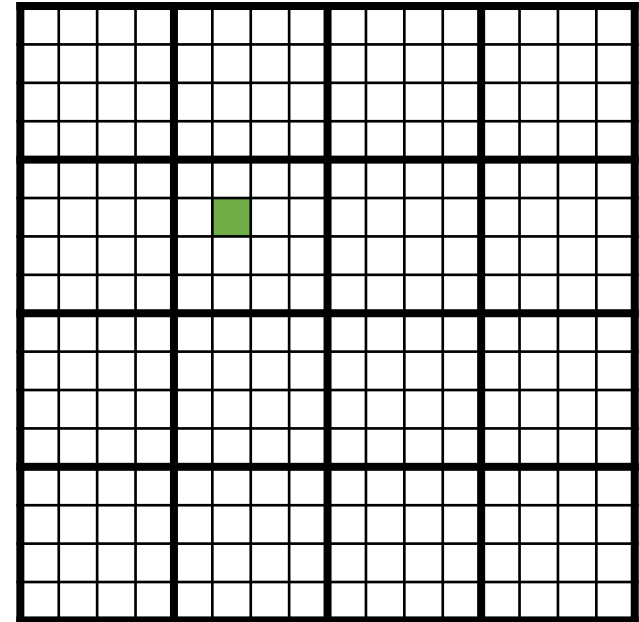
- Convolution
- **Parallel Convolution: a basic algorithm**
- Constant Memory and Caching
- Tiled Convolution



input




filter



output

Parallelization approach: Assign one thread to compute each **output element** by looping over **input elements** and **filter** weights

- Based on our experience in image smoothing and matrix multiplication, we can quickly write a **simple parallel convolution kernel**.



N: a pointer to the input array
F: a pointer to the filter
P: a pointer to the output array
r: the radius of the square filter
width: the width of the input and output arrays
height: the height of the input and output arrays


```
01 __global__ void convolution_2D_basic_kernel(float *N, float *F, float *P,
02     int r, int width, int height) {
03     int outCol = blockIdx.x*blockDim.x + threadIdx.x;
04     int outRow = blockIdx.y*blockDim.y + threadIdx.y;
05     float Pvalue = 0.0f;
06     for (int fRow = 0; fRow < 2*r+1; fRow++) {
07         for (int fCol = 0; fCol < 2*r+1; fCol++) {
08             inRow = outRow - r + fRow;
09             inCol = outCol - r + fCol;
10             if (inRow >= 0 && inRow < height && inCol >= 0 && inCol < width) {
11                 Pvalue += F[fRow][fCol]*N[inRow*width + inCol];
12             }
13         }
14     }
15     P[outRow][outCol] = Pvalue;
```

- Two concerns, regarding performance considerations,
 - Control Divergence**
 - The threads that calculate the output elements near the four edges of the P array will need to handle ghost cells. Therefore, **control divergence occurs in the if-statement**.
 - However, **for large input arrays and small filters**, control divergence occurs only in computing a small portion of the output elements, which will **keep the effect of control divergence small**.
 - Memory Bandwidth**
 - The compute to global memory access ratio is only about 0.25 OP/B
 - 2 operations for every 8 bytes loaded (on line 10).
 - We will need to **reduce the number of global memory accesses**.

```

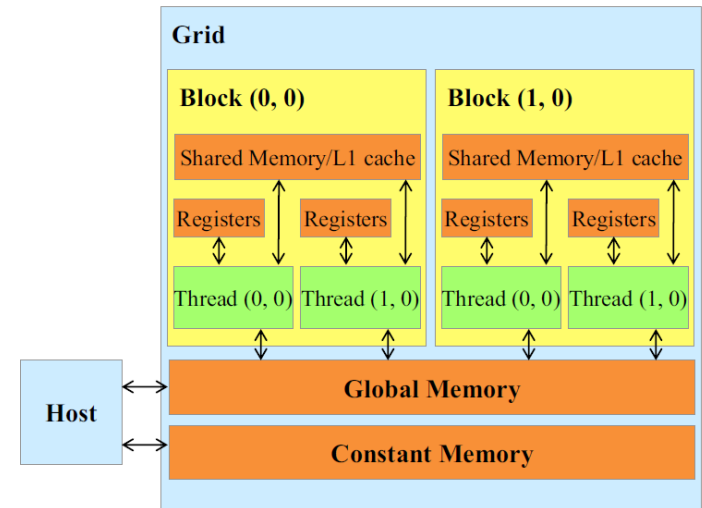

01 __global__ void convolution_2D_basic_kernel(float *N, float *F, float *P,
02   int r, int width, int height) {
03   int outCol = blockIdx.x*blockDim.x + threadIdx.x;
04   int outRow = blockIdx.y*blockDim.y + threadIdx.y;
05   float Pvalue = 0.0f;
06   for (int fRow = 0; fRow < 2*r+1; fRow++) {
07     for (int fCol = 0; fCol < 2*r+1; fCol++) {
08       inRow = outRow - r + fRow;
09       inCol = outCol - r + fCol;
10       if (inRow >= 0 && inRow < height && inCol >= 0 && inCol < width) {
11         Pvalue += F[fRow][fCol]*N[inRow*width + inCol];
12       }
13     }
14   }
15   P[outRow][outCol] = Pvalue;


```



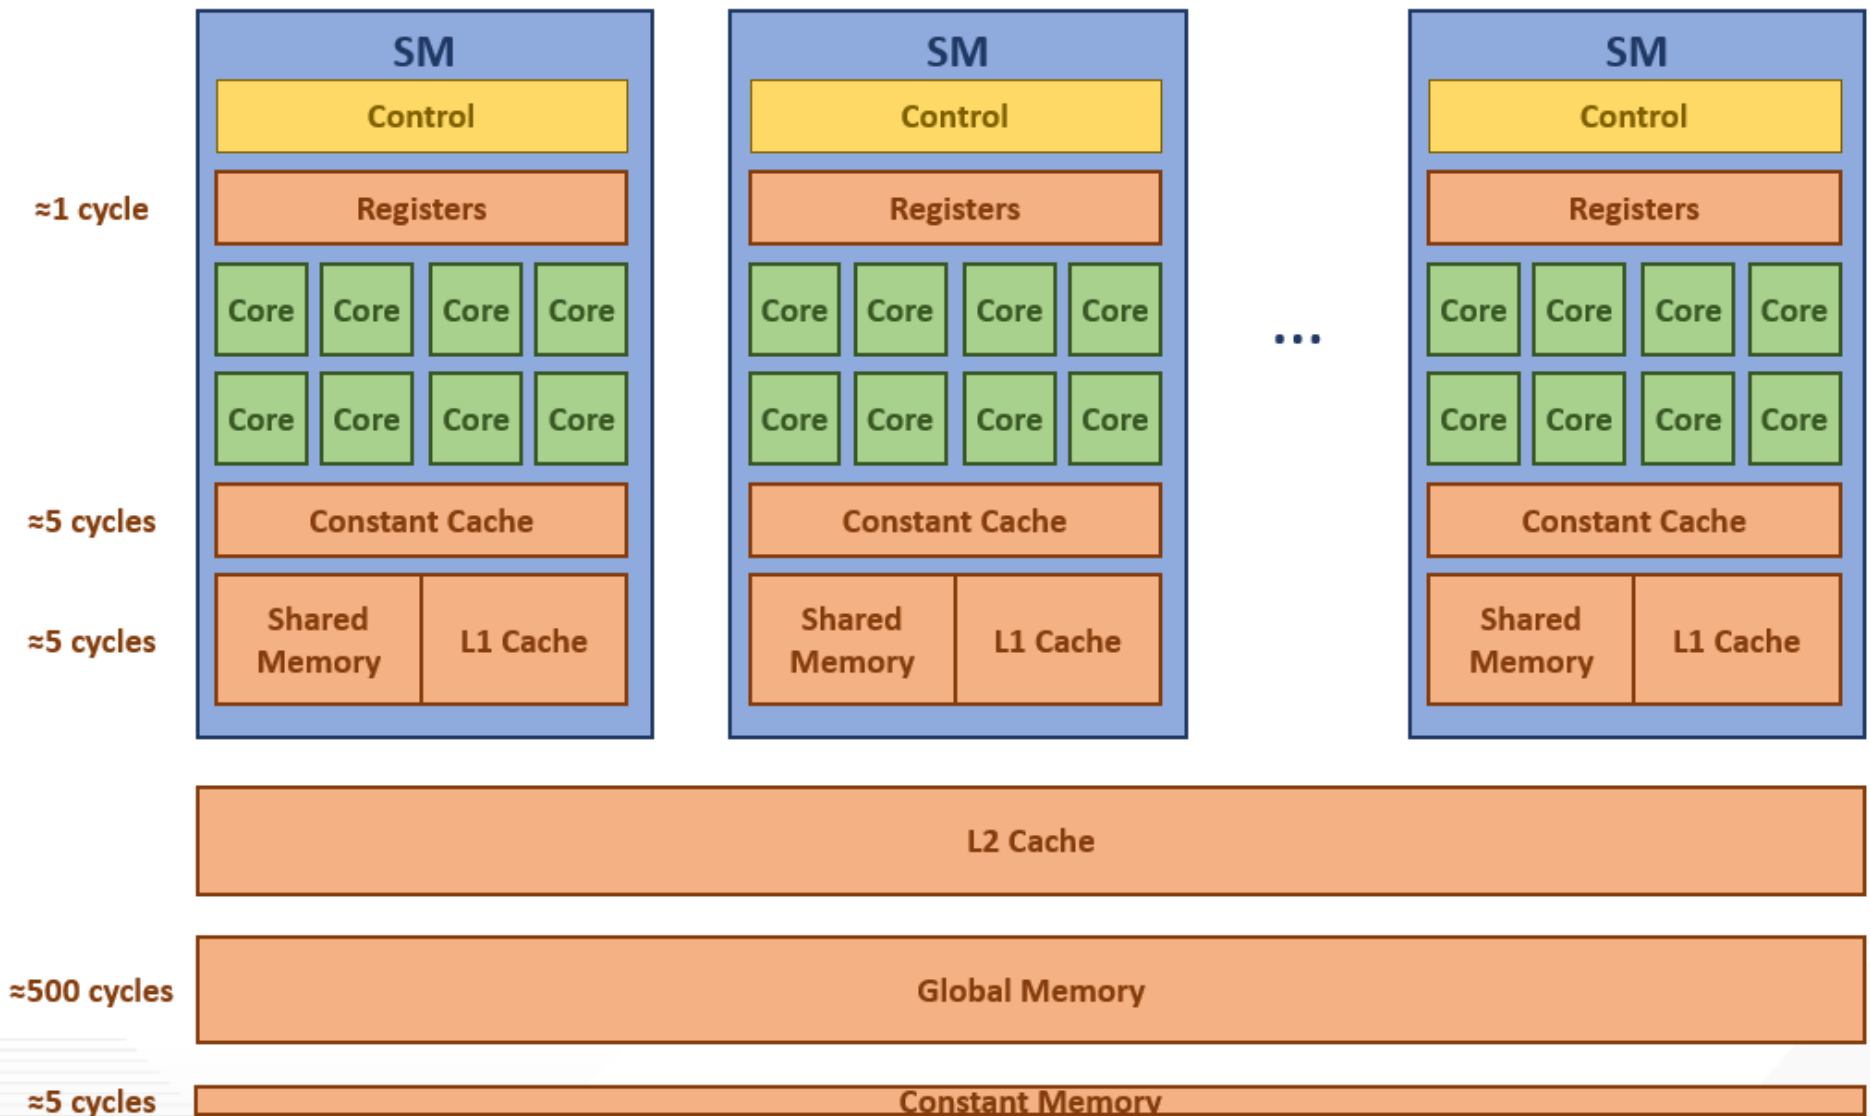
- Convolution
- Parallel Convolution: a basic algorithm
- **Constant Memory and Caching**
- Tiled Convolution

- Observations:
 - The filter is typically small
 - The filter is constant (weights do not change)
 - The filter is accessed by all threads in the grid
- These three properties make the filter an excellent candidate for **constant memory and caching**.
- Optimization: store the filter in **constant memory** for quicker access





Recall: Memory in the CUDA Programming Model



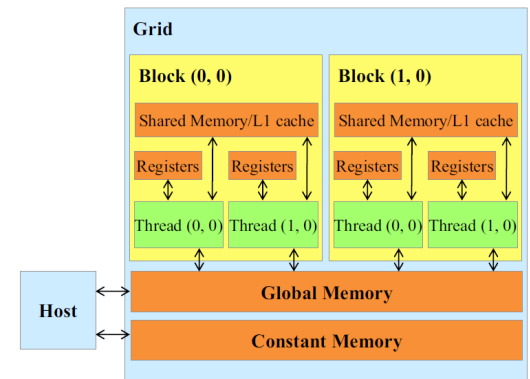

- Declare constant memory array as global variable

```
#define FILTER_RADIUS 2
__constant__ float F[2*FILTER_RADIUS+1][2*FILTER_RADIUS+1];
```

- Must initialize constant memory from the host:
 - Cannot modify during execution

```
cudaMemcpyToSymbol(F, F_h, (2*FILTER_RADIUS+1)*(2*FILTER_RADIUS+1)*sizeof(float));
```

- Can only allocate up to 64KB
 - Otherwise, input is also constant, but it is too large to put in constant memory




- A 2D convolution kernel **using constant memory for the filter F**
 - Note: Kernel functions access constant memory variables as global variables. Therefore, their pointers do not need to be passed to the kernel as arguments.


```

01 __global__ void convolution_2D_const_mem_kernel(float *N, float *P, int r,
    int width, int height) {
02     int outCol = blockIdx.x*blockDim.x + threadIdx.x;
03     int outRow = blockIdx.y*blockDim.y + threadIdx.y;
04     float Pvalue = 0.0f;
05     for (int fRow = 0; fRow < 2*r+1; fRow++) {
06         for (int fCol = 0; fCol < 2*r+1; fCol++) {
07             inRow = outRow - r + fRow;
08             inCol = outCol - r + fCol;
09             if (inRow >= 0 && inRow < height && inCol >= 0 && inCol < width) {
10                 Pvalue += F[fRow][fCol]*N[inRow*width + inCol];
11             }
12         }
13     }
14     P[outRow*width+outCol] = Pvalue;
15 }


```



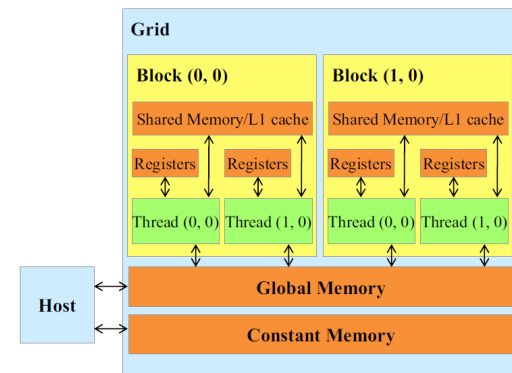
input



filter



output

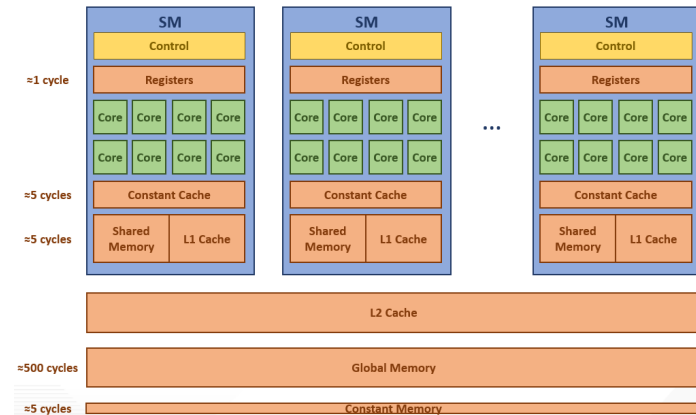


- Constant data: easier to build an efficient cache
 - No need to support write back
- Constant Cache
 - A small, specialized cache can be highly effective in capturing the heavily used constant memory variable for each kernel.
- (Slightly Improved) Memory Bandwidth
 - The compute to global memory access ratio is now about 0.5 OP/B
 - 2 operations for every 4 bytes loaded (on line 10).
 - However, **further optimization is needed to reduce the number of global memory accesses.**


```

01 __global__ void convolution_2D_const_mem_kernel(float *N, float *P, int r,
02     int width, int height) {
03     int outCol = blockIdx.x*blockDim.x + threadIdx.x;
04     int outRow = blockIdx.y*blockDim.y + threadIdx.y;
05     float Pvalue = 0.0f;
06     for (int fRow = 0; fRow < 2*r+1; fRow++) {
07         for (int fCol = 0; fCol < 2*r+1; fCol++) {
08             inRow = outRow - r + fRow;
09             inCol = outCol - r + fCol;
10             if (inRow >= 0 && inRow < height && inCol >= 0 && inCol < width) {
11                 Pvalue += F[fRow][fCol]*N[inRow*width + inCol];
12             }
13         }
14     }
15     P[outRow*width+outCol] = Pvalue;


```




- Convolution
- Parallel Convolution: a basic algorithm
- Constant Memory and Caching
- **Tiled Convolution**



input
(in global
memory)




filter
(in constant
memory)




output


Observation: Threads in the same block load some of the same input elements



input
(in global
memory)




filter
(in constant
memory)




output


Observation: Threads in the same block load some of the same input elements



input
(in global
memory)




filter
(in constant
memory)




output

Observation: Threads in the same block load some of the same input elements


Optimization: Each thread loads one input element to shared memory and other threads access the element from shared memory



$\text{input}_{\text{tile}}$
(in shared
memory)




filter
(in constant
memory)



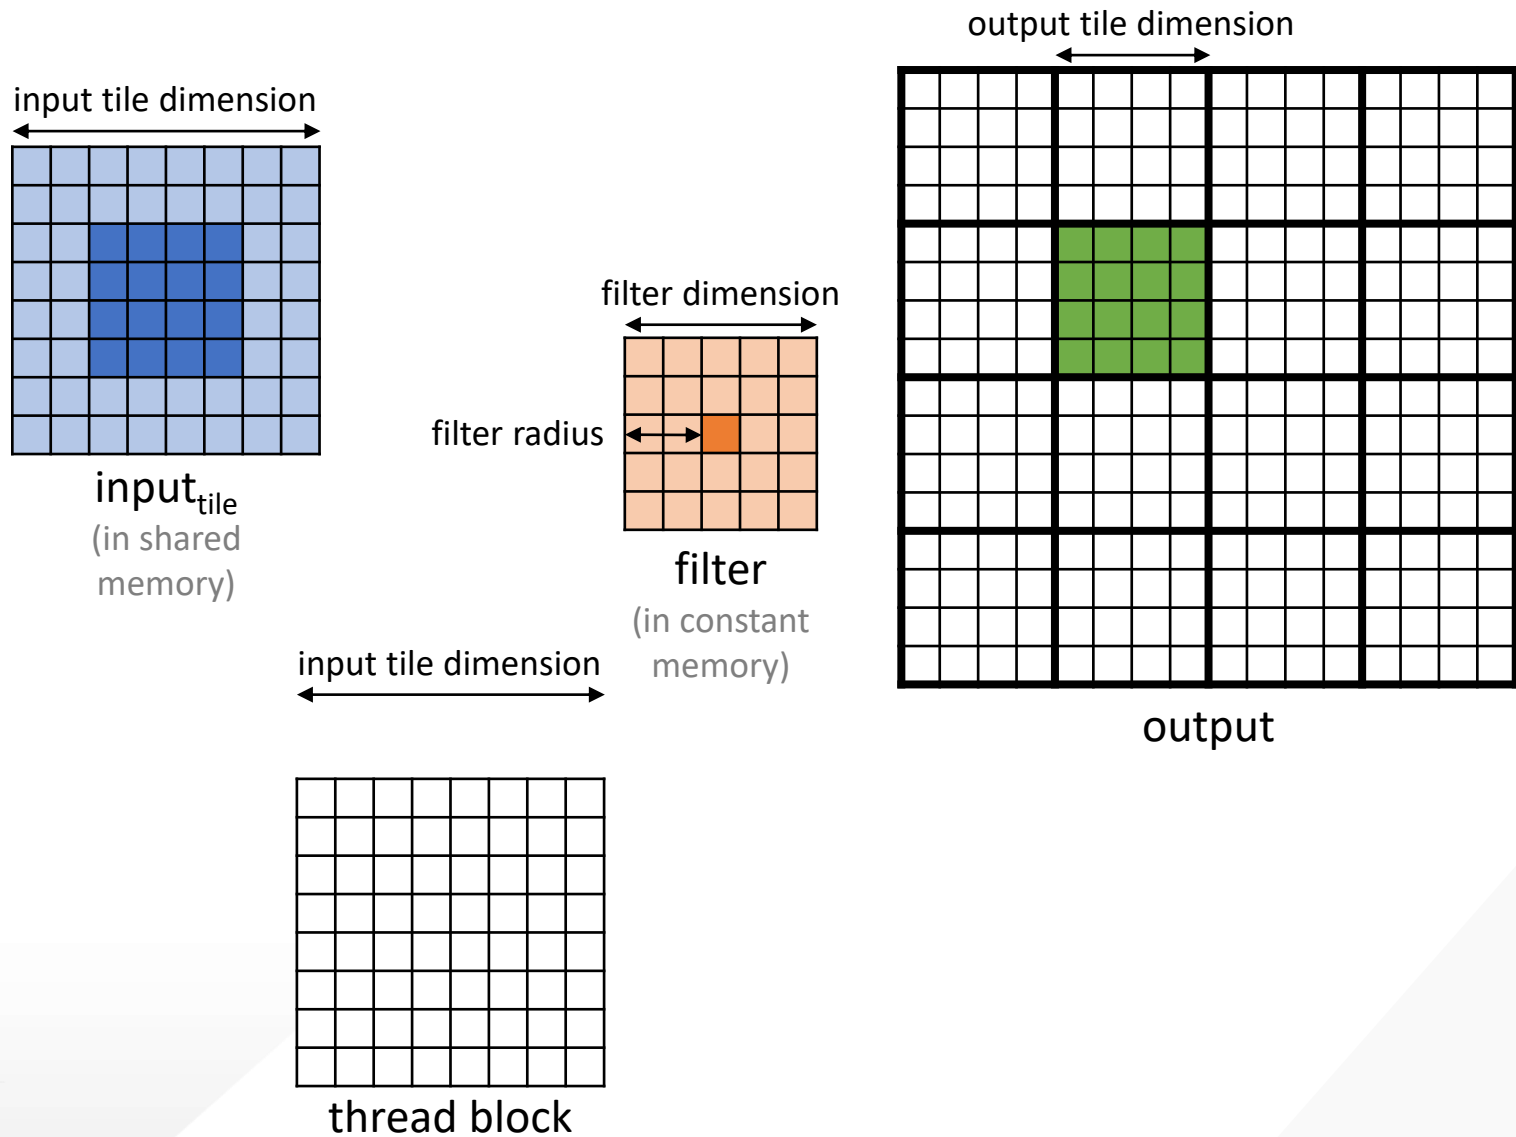
output

Optimization: Each thread loads one input element to shared memory and other threads access the element from shared memory




Challenge: Input and output tiles have different dimensions
 ($\text{input tile dimension} = \text{output tile dimension} + 2 \times \text{filter radius}$)

Solution: Launch enough threads per block to load the input tile to shared memory, then use a subset of them to compute and store the output tile




input tile dimension



$\text{input}_{\text{tile}}$
(in shared
memory)

filter dimension




filter radius

filter


(in constant
memory)

input tile dimension



thread block


output tile dimension



output


threads active when
loading input tile

input tile dimension



$\text{input}_{\text{tile}}$
(in shared
memory)


filter dimension



filter radius


filter
(in constant
memory)

input tile dimension
output tile dimension



thread block

output tile dimension



output


threads active when
computing and storing
the output tile

- A tiled 2D convolution kernel

```

01 #define IN_TILE_DIM 32
02 #define OUT_TILE_DIM ((IN_TILE_DIM) - 2*(FILTER_RADIUS))
03 __constant__ float F[2*FILTER_RADIUS+1][2*FILTER_RADIUS+1];
04 __global__ void convolution_tiled_2D_const_mem_kernel(float *N, float *P,
05                                                       int width, int height) {
06     int col = blockIdx.x*OUT_TILE_DIM + threadIdx.x - FILTER_RADIUS;
07     int row = blockIdx.y*OUT_TILE_DIM + threadIdx.y - FILTER_RADIUS;
08     //loading input tile
09     shared N_s[IN_TILE_DIM][IN_TILE_DIM];
10     if(row>=0 && row<height && col>=0 && col<width) {
11         N_s[threadIdx.y][threadIdx.x] = N[row*width + col];
12     } else {
13         N_s[threadIdx.y][threadIdx.x] = 0.0;
14     }
15     syncthreads();
16     // Calculating output elements
17     int tileCol = threadIdx.x - FILTER_RADIUS;
18     int tileRow = threadIdx.y - FILTER_RADIUS;
19     // turning off the threads at the edges of the block
20     if (col >= 0 && col < width && row >= 0 && row < height) {
21         if (tileCol>=0 && tileCol<OUT_TILE_DIM && tileRow>=0
22             && tileRow<OUT_TILE_DIM){
23             float Pvalue = 0.0f;
24             for (int fRow = 0; fRow < 2*FILTER_RADIUS+1; fRow++) {
25                 for (int fCol = 0; fCol < 2*FILTER_RADIUS+1; fCol++) {
26                     Pvalue += F[fRow][fCol]*N_s[tileRow+fRow][tileCol+fCol];
27                 }
28             }
29             P[row*width+col] = Pvalue;
30         }
31     }
32 }

```





- A tiled 2D convolution kernel

```

01 #define IN_TILE_DIM 32
02 #define OUT_TILE_DIM ((IN_TILE_DIM) - 2*(FILTER_RADIUS))
03 __constant__ float F[2*FILTER_RADIUS+1][2*FILTER_RADIUS+1];
04 __global__ void convolution_tiled_2D_const_mem_kernel(float *N, float *P,
05                                                     int width, int height) {
06     int col = blockIdx.x*OUT_TILE_DIM + threadIdx.x - FILTER_RADIUS;
07     int row = blockIdx.y*OUT_TILE_DIM + threadIdx.y - FILTER_RADIUS;
08     //loading input tile
09     shared N_s[IN_TILE_DIM][IN_TILE_DIM];
10     if(row>=0 && row<height && col>=0 && col<width) {
11         N_s[threadIdx.y][threadIdx.x] = N[row*width + col];
12     } else {
13         N_s[threadIdx.y][threadIdx.x] = 0.0;
14     }
15     syncthreads();
16     // Calculating output elements
17     int tileCol = threadIdx.x - FILTER_RADIUS;
18     int tileRow = threadIdx.y - FILTER_RADIUS;
19     // turning off the threads at the edges of the block
20     if (col >= 0 && col < width && row >= 0 && row < height) {
21         if (tileCol>=0 && tileCol<OUT_TILE_DIM && tileRow>=0
22             && tileRow<OUT_TILE_DIM) {
23             float Pvalue = 0.0f;
24             for (int fRow = 0; fRow < 2*FILTER_RADIUS+1; fRow++) {
25                 for (int fCol = 0; fCol < 2*FILTER_RADIUS+1; fCol++) {
26                     Pvalue += F[fRow][fCol]*N_s[tileRow+fRow][tileCol+fCol];
27                 }
28             }
29             P[row*width+col] = Pvalue;
30         }
31     }
32 }

```

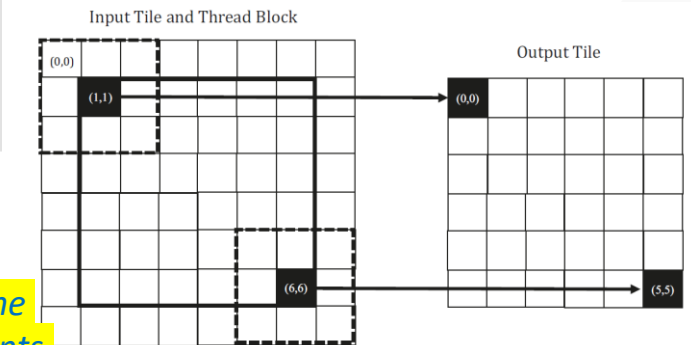
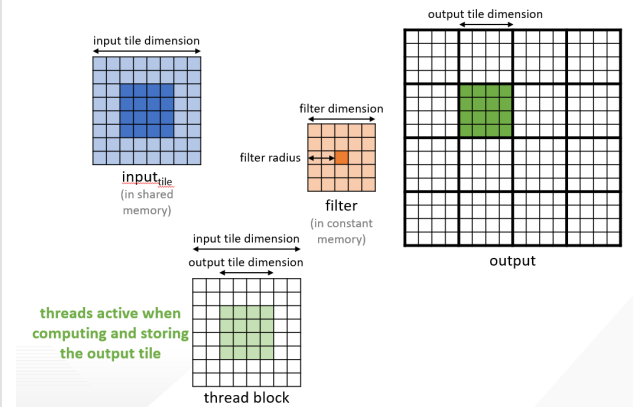



- A tiled 2D convolution kernel

```

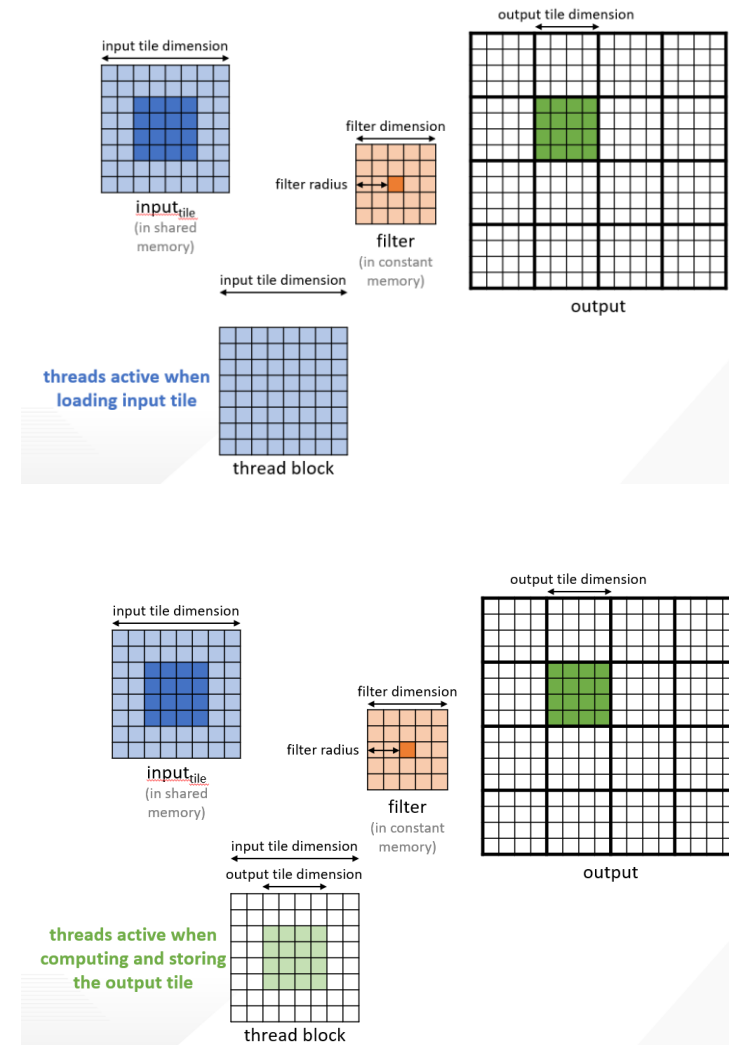
01 #define IN_TILE_DIM 32
02 #define OUT_TILE_DIM ((IN_TILE_DIM) - 2*(FILTER_RADIUS))
03 __constant__ float F[2*FILTER_RADIUS+1][2*FILTER_RADIUS+1];
04 __global__ void convolution_tiled_2D_const_mem_kernel(float *N, float *P,
05                                                     int width, int height) {
06     int col = blockIdx.x*OUT_TILE_DIM + threadIdx.x - FILTER_RADIUS;
07     int row = blockIdx.y*OUT_TILE_DIM + threadIdx.y - FILTER_RADIUS;
08     //loading input tile
09     shared N_s[IN_TILE_DIM][IN_TILE_DIM];
10     if(row>=0 && row<height && col>=0 && col<width) {
11         N_s[threadIdx.y][threadIdx.x] = N[row*width + col];
12     } else {
13         N_s[threadIdx.y][threadIdx.x] = 0.0;
14     }
15     syncthreads();
16     // Calculating output elements
17     int tileCol = threadIdx.x - FILTER_RADIUS;
18     int tileRow = threadIdx.y - FILTER_RADIUS;
19     // turning off the threads at the edges of the block
20     if (col >= 0 && col < width && row >= 0 && row < height) {
21         if (tileCol>=0 && tileCol<OUT_TILE_DIM && tileRow>=0
22             && tileRow<OUT_TILE_DIM) {
23             float Pvalue = 0.0f;
24             for (int fRow = 0; fRow < 2*FILTER_RADIUS+1; fRow++) {
25                 for (int fCol = 0; fCol < 2*FILTER_RADIUS+1; fCol++) {
26                     Pvalue += F[fRow][fCol]*N_s[tileRow+fRow][tileCol+fCol];
27                 }
28             }
29             P[row*width+col] = Pvalue;
30         }
31     }
32 }


```



A small example that illustrates the thread organization for using the input tile elements in the shared memory to calculate the output tile elements.

- With tiling:
 - Considering output tile dimension:
 - Input = IN_TILE_DIM
 - Output = OUT_TILE_DIM
 - Global loads per block: $IN_TILE_DIM^2 * 4 = (OUT_TILE_DIM + 2 * FILTER_RADIUS)^2 * 4$
 - Each thread that is assigned to an input tile element loads one 4-byte input value.
 - Operations per block: $OUT_TILE_DIM^2 * (2 * FILTER_RADIUS + 1)^2 * 2$
 - Every thread that is assigned to an output tile element, performs 1 multiplication and 1 addition for every element of the filter
 - Ratio:
$$\frac{OUT_TILE_DIM^2 * (2 * FILTER_RADIUS + 1)^2 * 2}{(OUT_TILE_DIM + 2 * FILTER_RADIUS)^2 * 4}$$
 - For example, when $FILTER_RADIUS = 2$ and $OUT_TILE_DIM = 28$, the ratio is 9.57 OP/B ($\approx 19\times$ improvement!)





Threads computing output elements at the boundary access input elements that are out of bounds (also called *ghost* elements)

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0						
0	0						
0	0						
0	0						
0	0						
0	0						

$\text{input}_{\text{tile}}$
(in shared
memory)

filter
(in constant
memory)

output

Threads computing output elements at the boundary access input elements that are out of bounds (also called *ghost* elements)

Solution: Store zero to shared memory tile for our of bounds input elements

- Wen-mei W. Hwu, David B. Kirk, and Izzat El Hajj. *Programming Massively Parallel Processors: A Hands-on Approach*. Morgan Kaufmann, 2022.